Computing the Lambert W function in arbitrary-precision complex interval arithmetic

Fredrik Johansson*

Abstract

We describe an algorithm to evaluate all the complex branches of the Lambert W function with rigorous error bounds in interval arithmetic, which has been implemented in the Arb library. The classic 1996 paper on the Lambert W function by Corless *et al.* provides a thorough but partly heuristic numerical analysis which needs to be complemented with some explicit inequalities and practical observations about managing precision and branch cuts.

1 Introduction

The Lambert W function $W(z)$ is the inverse function of $f(w) = we^w$, meaning that $W(z)e^{W(z)} = z$ holds for any z. Since f is not injective, the Lambert W function is multivalued, having an infinite number of branches $W_k(z)$, $k \in \mathbb{Z}$, analogous to the branches $\ln_k(z) = \log(z) + 2\pi i k$ of the natural logarithm which inverts $g(w) = e^w$.

The study of the equation $we^w = z$ goes back to Lambert and Euler in the 18th century, but a standardized notation for the solution only appeared in the 1990s with the introduction of `LambertW` in the Maple computer algebra system, along with the paper [2] by Corless, Gonnet, Hare, Jeffrey and Knuth which collected and proved the function's main properties. There is now a vast literature on applications, and in 2016 a conference was held to celebrate the first 20 years of the Lambert W function.

The paper [2] sketches how $W_k(z)$ can be computed for any $z \in \mathbb{C}$ and any k, using a combination of series expansions and iterative root-finding. Numerical implementations are available in many computer algebra systems and numerical libraries; see for instance [6, 1, 8]. However, there is no published work to date addressing interval arithmetic or discussing a complete rigorous implementation of the complex branches.

The equation $we^w - z = 0$ can naturally be solved with any standard interval root-finding method like subdivision or the interval Newton method [7]. Another possibility, suggested in [2], is to use a posteriori error analysis to bound the error of an approximate solution. The Lambert W function can also be evaluated as the solution of an ordinary differential equation, for which rigorous solvers are available. Regardless of the approach, the main difficulty is to make sure that correctness and efficiency are maintained near singularities and branch cuts.

This paper describes an algorithm for rigorous evaluation of the Lambert W function in complex interval arithmetic, which has been implemented in the Arb library [4]. This implementation was designed to achieve the following goals:

- $W(z)$ is only a constant factor more expensive to compute than elementary functions like $\log(z)$ or $\exp(z)$. For rapid, rigorous computation of elementary functions in arbitrary precision, the methods in [3] are used.

*LFANT project-team, INRIA Bordeaux-Sud-Ouest. Contact: `fredrik.johansson@gmail.com`.

Figure 1: Plot of the real branches $W_0(x)$ and $W_{-1}(x)$ computed with Arb. The boxes show the size of the output intervals given wide input intervals. In this plot, the input intervals have been subdivided until the output radius is smaller than ε.

- The output enclosures are reasonably tight.
- All the complex branches W_k are supported, with a stringent treatment of branch cuts.
- It is possible to compute derivatives $W^{(n)}(z)$ efficiently, for arbitrary n.

The main contribution of this paper is to derive bounds with explicit constants for a posteriori certification and for the truncation error in certain series expansions, in cases where previous publications give big-O estimates. We also discuss the implementation of the complex branches in detail.

Arb uses (extended) real intervals of the form $[m \pm r]$, shorthand for $[m - r, m + r]$, where the midpoint m is an arbitary-precision floating-point number and the radius r is an unsigned fixed-precision floating-point number. The exponents of m and r are bignums which can be arbitrarily large (this is useful for asymptotic problems, and removes edge cases with underflow or overflow). Complex numbers are represented in rectangular form $x + yi$ using pairs of real intervals. We will occasionally rely on these implementation details, but generally speaking the methods translate easily to other interval formats.

1.1 Complex branches

In this work, $W_k(z)$ always refers to the standard k-th branch as defined in [2]. We sometimes write $W(z)$ when referring to the multivalued Lambert W function or a branch implied by the context. Before we proceed, we summarize the branch structure of W. A more detailed description with illustrations can be found in [2].

Figure 1 demonstrates evaluation of the Lambert W function in the two real-valued regions. The *principal branch* $W_0(z)$ is real-valued and monotone increasing for real $z \geq -1/e$, with the image $[-1, \infty)$, while $W_{-1}(z)$ is real-valued and monotone decreasing for real $-1/e \leq z < 0$, with the image $(-\infty, -1]$. Everywhere else, $W_k(z)$ is complex. There is a square root-type singularity at the branch point $z = -1/e$ connecting the real segments, where $W_0(-1/e) = W_{-1}(-1/e) = -1$. The principal branch contains the root $W_0(0) = 0$,

which is the only root of W. For all $k \neq 0$, the point $z = 0$ is a branch point with a logarithmic singularity.

$W_0(z)$ has a single branch cut on $(-\infty, -1/e)$, while the branches $W_k(z)$ with $|k| \geq 2$ have a single branch cut on $(-\infty, 0)$. The branches $W_{\pm 1}$ are more complicated, with a set of adjacent branch cuts: in the upper half plane, W_{-1} has a branch cut on $(-\infty, -1/e)$ and one on $(-1/e, 0)$; in the lower half plane, W_{-1} has a single branch cut on $(-\infty, 0)$. W_1 is similar to W_{-1}, but with the sides exchanged. The branch cuts on $(-\infty, 0)$ or $(-\infty, -1/e)$ connect W_k with W_{k+1}, while the branch cuts on $(-1/e, 0)$ connect W_{-1} with W_1.

We follow the convention that the function value on a branch cut is continuous when approaching the cut in the counterclockwise direction around a branch point. For the standard branches $W_k(z)$, this is the same as continuity with the upper half plane, i.e. $W_k(x + 0i) = \lim_{y \to 0+} W_k(x + yi)$. When $\mathrm{Im}(z) \neq 0$, we have $W_k(z) = \overline{W_{-k}(\bar{z})}$. By the same convention, the principal branch of the natural logarithm is defined to satisfy $\mathrm{Im}(\log(z)) \in (-\pi, +\pi]$.

We do not use signed zero in the sense of IEEE 754 floating-point arithmetic, which would allow preserving continuity from either side of a branch cut. This is a trivial omission since we can distinguish between $W(x + 0i)$ and $W(x - 0i)$ using $W_k(x - 0i) = \overline{W_{-k}(x + 0i)}$.

In interval arithmetic, we need to enclose the union of the images of $W(z)$ on both sides of the cut when the interval representing z straddles a branch cut. The jump discontinuity between the cuts will prevent the output intervals from converging when the input intervals shrink (unless the input intervals lie exactly on a branch cut, say $z = [-5, -4] + 0i$). This problem is solved by providing a set of alternative branch cuts to complement the standard cuts, as discussed in Section 4.

2   The main algorithm

The algorithm to evaluate the Lambert W function has three main ingredients:

- (Asymptotic cases.) If $|z|$ is extremely small or large, or if z is extremely close to the branch point at $-1/e$ when $W(z) \approx -1$, use the respective Taylor, Puiseux or asymptotic series to compute $W(z)$ directly.

- (Approximation.) Use floating-point arithmetic to compute some $\tilde{w} \approx W(\mathrm{mid}(z))$.

- (Certification.) Given $\tilde{w}$, use interval arithmetic (or floating-point arithmetic with directed rounding) to determine a bound r such that $|W(z) - \tilde{w}| \leq r$, and return $\tilde{w} + [\pm r] + [\pm r]i$, or simply $[\tilde{w} \pm r]$ when $W(z)$ is real-valued.

The special treatment of asymptotic cases is not necessary, but improves performance since the error can be bounded directly without a separate certification step. We give error bounds for the truncated series expansions in Section 3.

Computing a floating-point approximation with heuristic error control is a well understood problem, and we avoid going into too much detail here. Essentially, Arb uses the Halley iteration

$$w_{j+1} = w_j - \frac{w_j e^{w_j} - \mathrm{mid}(z)}{e^{w_j+1} - \dfrac{(w_j + 2)(w_j e^{w_j} - \mathrm{mid}(z))}{2w_j + 2}}$$

suggested in [2] to solve $we^w - \mathrm{mid}(z) = 0$, starting from a well-chosen initial value. In the most common cases, machine `double` arithmetic is first used to achieve near 53-bit accuracy (with care to avoid overflow or underflow problems or loss of significance near $z = -1/e$).

For typical accuracy goals of less than a few hundred bits, this leaves at most a couple of iterations to be done using arbitrary-precision arithmetic.

In the arbitrary-precision phase, the working precision is initially set low and then increases with each Halley iteration step to match the estimated number of accurate bits (which roughly triples with each iteration). This ensures that obtaining p accurate bits costs $O(1)$ full-precision exponential function evaluations instead of $O(\log p)$.

2.1 Certification

To compute a certified error bound for $\tilde{w}$, we use backward error analysis, following the suggestion of [2]. We compute $\tilde{z} = \tilde{w}e^{\tilde{w}}$ with interval arithmetic, and use

$$\tilde{w} = W(\tilde{z}) = W(z) + \int_z^{\tilde{z}} W'(t)dt. \tag{1}$$

to bound the error $W_k(\tilde{z}) - W_k(z)$. This approach relies on having a way to bound $|W_k'|$, which we address in Section 3.

The formal identity (1) is only valid provided that the correct integration path is taken on the Riemann surface of the multivalued W function. During the certification, we verify that the straight-line path γ from z to $\tilde{z}$ for W_k is correct in (1), so that the error is bounded by $|z - \tilde{z}| \sup_{t \in \gamma} |W_k'(t)|$. This is essentially to say that we have approximated $W_k(z)$ for the right k, since a poor starting value (or rounding error) in the Halley iteration could have put $\tilde{w}$ on the wrong branch, or closer to a solution on the wrong branch than the intended solution.

Algorithm 1 Compute certified enclosure of $W_k(z)$. The input is a complex interval z, a branch index $k \in \mathbb{Z}$, and a complex floating-point number $\tilde{w}$.

---

1. Verify that $\tilde{w} = x + yi$ lies in the range of the branch W_k:

   (a) Compute $t = x\operatorname{sinc}(y)$, $v = -\cos(y)$, $u = \operatorname{sgn}(k)y/\pi$ using interval arithmetic.
   (b) If $k = 0$, check $(|u| < 1) \wedge (t > v)$.
   (c) If $k \neq 0$, check $P_1 \wedge (P_2 \vee P_3 \vee P_4)$ where

   $$\begin{aligned}
   P_1 &= (u > 2|k| - 2) \wedge (u < 2|k| + 1) \\
   P_2 &= (u > 2|k| - 1) \wedge (u < 2|k|) \\
   P_3 &= (u < 2|k|) \wedge (t < v) \\
   P_4 &= (u > 2|k| - 1) \wedge (t > v).
   \end{aligned}$$

   (d) If the check fails, return $[\pm\infty] + [\pm\infty]i$.

2. Compute $\tilde{z} = \tilde{w}e^{\tilde{w}}$ using interval arithmetic.
3. Compute a complex interval $U \supseteq z \cup \tilde{z}$ (U will contain the straight line from z to $\tilde{z}$).
4. Verify that U does not cross a branch cut: check

   $$(\operatorname{Im}(U) \geq 0) \vee (\operatorname{Im}(U) < 0) \vee \begin{pmatrix} \operatorname{Re}(eU + 1) > 0 & \text{if } k = 0 \\ \operatorname{Re}(U) > 0 & \text{if } k \neq 0 \end{pmatrix}.$$

   If the check fails, return $[\pm\infty] + [\pm\infty]i$.
5. Compute a bound $C \geq |W_k'(U)|$ and return $\tilde{w} + [\pm r] + [\pm r]i$ where $r = C|z - \tilde{z}|$.

The complete certification procedure is stated in Algorithm 1. In the pseudocode, all pointwise predicates are extended to intervals in the strong sense; for example, $x \geq 0$ evalu-

ates to true if all points in the interval representing x are nonnegative, and false otherwise. A predicate that should be true for exact input in infinite precision arithmetic can therefore evaluate to false due to interval overestimation or insufficient precision.

In the first step, we use the fact that the images of the branches in the complex W-plane are separated by the line $(-\infty, -1/e]$ together with the curves $\{-\eta \cot \eta + \eta i\}$ for $-\pi < \eta < \pi$ and $2k\pi < \pm\eta < (2k+1)\pi$ (this is proved in [2]). In the $k \neq 0$ case, the predicates P_2, P_3, P_4 cover overlapping regions, allowing the test to pass even if $\tilde{w}$ falls very close to one of the curves with $2k\pi < \pm\eta < (2k+1)\pi$ where a sign change occurs, i.e. when z crosses the real axis to the right of the branch point.

The test in Algorithm 1 always fails when z lies on a branch cut, or too close to a cut to resolve with a reasonable precision, say if $z = -2^{10^{10}} + 10i$ or $z = -10 + 2^{-10^{10}}i$. This problem could be solved by taking the location of z into account in addition that of $\tilde{w}$. In Arb, a different solution has been implemented, namely to perturb z away from the branch cut before calling Algorithm 1 (together with an error bound for this perturbation). This works well in practice with the use of a few guard bits, and seemed to require less extra logic to implement.

Due to the cancellation in evaluating the residual $z - \tilde{z}$, the quantity $\tilde{z} = \tilde{w}e^{\tilde{w}}$ needs to be computed to at least p-bit precision in the certification step to achieve a relative error bound of 2^{-p}. Here, a useful optimization is to compute e^{w_j} with interval arithmetic in the last Halley update $\tilde{w} = w_{j+1} = H(w_j)$ and then compute $e^{\tilde{w}}$ as $e^{w_j}e^{\tilde{w}-w_j}$. Evaluating $e^{\tilde{w}-w_j}$ costs only a few series terms of the exponential function since $|\tilde{w} - w_j| \approx 2^{-p/3}$.

A different possibility for the certification step would be to guess an interval around $\tilde{w}$ and perform one iteration with the interval Newton method. This can be combined with the main iteration, simultaneously extending the accuracy from $p/2$ to p bits and certifying the error bound. An advantage of the interval Newton method is that it operates directly on the function $f(w) = we^w - z$ and its derivative without requiring explicit knowledge about W'. This method was tested but ultimately abandoned in the Arb implementation since it seemed more difficult to handle the precision and make a good interval guess in practice, particularly when z is represented by a wide interval. In any case the branch certification would still be necessary.

2.2   The main algorithm in more detail

Algorithm 2 describes the main steps implemented by the Arb function with signature

```
void acb_lambertw(acb_t res, const acb_t z,
   const fmpz_t k, int flags, slong prec)
```

where `acb_t` denotes Arb's complex interval type, `res` is the output variable, `fmpz_t` is a multiprecision integer type, and `prec` gives the precision goal p in bits.

In step 2, we switch to separate code for real-valued input and output (calling the function `arb_lambertw` which uses real `arb_t` interval variables). The real version implements essentially the same algorithm as the complex version, but skips most branch cut related logic.

In step 3, we reduce the working precision to save time if the input is known to less than p accurate bits. The precision is subsequently adjusted in step 5, accounting for the fact that we gain accurate bits in the value of $W_k(z)$ from the exponent of mid(z) or k when $|W_k(z)|$ is large. Step 5 is cheap, as it only requires inspecting the exponents of the floating-point components of z and computing bit lengths of integers.

The constants T, L, M, P appearing in steps 4, 6 and 7 are tuning parameters to control the number of series expansion terms allowed to compute W directly instead of falling back to

Algorithm 2 Main algorithm for $W_k(z)$ implemented in `acb_lambertw`. The input is a complex interval z, a branch index $k \in \mathbb{Z}$, and a precision $p \in \mathbb{Z}_{\geq 2}$.

1. If z is not finite or if $k \neq 0$ and $0 \in z$, return indeterminate ($[\pm\infty] + [\pm\infty]i$).

2. If $k = 0$ and $z \subset (-1/e, \infty)$, or if $k = -1$ and $z \subset (-1/e, 0)$, return $W_k(z)$ computed using dedicated code for the real branches.

3. Set the accuracy goal to $q \leftarrow \min(p, \max(10, -\log_2 \operatorname{rad}(z)/|\operatorname{mid}(z)|))$.

4. If $k = 0$ and $|\operatorname{mid}(z)| < 2^{-q/T}$, return $W_0(z)$ computed using T terms of the Taylor series.

5. Compute positive integers $b_1 \approx \log_2(|\log(z) + 2\pi ik|)$, $b_2 \approx \log_2(b_1)$. If $|z|$ is near ∞, or near 0 and $k \neq 0$, adjust the goal to $q \leftarrow \min(p, \max(q + b_1 - b_2, 10))$.

6. Let $s = 2 - b_1$, $t = 2 + b_2 - b_1$. If $b_1 - \max(t + Ls, Mt) > q$, return $W_k(z)$ computed using the asymptotic series with (L, M) terms.

7. Check if z is near the branch point at $-1/e$: if $|ez + 1| < 2^{-2q/P-6}$, and $|k| \leq 1$ (and $\operatorname{Im}(z) < 0$ if $k = 1$, or $\operatorname{Im}(z) \geq 0$ if $k = -1$) return $W_k(z)$ computed using P terms of the Puiseux series.

8. If z contains points on both sides of a branch cut, set $z_a = \operatorname{Re}(z) + (\operatorname{Im}(z) \cap [0, \infty))i$ and $z_b = \operatorname{Re}(z) + (-\operatorname{Im}(z) \cap [0, \infty))i$. Then compute $w_a = W_k(z_a)$ and $w_b = \overline{W_{-k}(z_b)}$ and return $w_a \cup w_b$.

9. Let $x + yi = \operatorname{mid}(z)$. If x lies to the left of a branch point (0 or $-1/e$) and $|y| < 2^{-q}|x|$, set $z' = \operatorname{Re}(z) + [\underline{\varepsilon \pm \varepsilon}]i$ where $\varepsilon = 2^{-q}|x|$ (if $y < 0$ in this case, modify the following steps to compute $\overline{W_{-k}(z')}$ instead of $W_k(z')$). Otherwise, set $z' = z$.

10. Compute a floating-point approximation $\tilde{w} \approx W_k(\operatorname{mid}(z'))$ to a heuristic accuracy of q bits plus a few guard bits.

11. Convert $\tilde{w}$ to a certified complex interval w for $W_k(\operatorname{mid}(z'))$ by calling Algorithm 1.

12. If z' is inexact, bound $|W_k'(z')| \leq C$ and add $[\pm r] + [\pm r]i$ to w, where $r = C \operatorname{rad}(z')$. Return w.

---

root-finding. These parameters could be made precision-dependent to optimize performance, but for most purposes small constants work well.

Step 8 ensures that z lies on one side of a branch cut, splitting the evaluation of $W_k(z)$ into two subcases if necessary. This step ensures that step 12 (which bounds the propagated error due to the uncertainty in z) is correct, since our bound for W' does not account for the branch cut jump discontinuity (and in any case differentiating a jump discontinuity would give the output $[\pm\infty] + [\pm\infty]i$ which is needlessly pessimistic). We note that conjugation is used to get a continuous evaluation of $W_k(\operatorname{Re}(z) + (\operatorname{Im}(z) \cap (-\infty, 0))i)$, in light of our convention to work with closed intervals and make the standard branches W_k continuous from above on the cut.

We perform step 8 *after* checking if the asymptotic series or Puiseux series can be used, since correctly implemented complex logarithm and square root functions take care of branch cuts automatically. If z needs to be split into z_a and z_b in step 8, then the main algorithm can be called recursively, but the first few steps can be skipped. However, step 7 should be repeated when $k = \pm 1$ since the Puiseux series near $-1/e$ might be valid for z_a or z_b even when it is not applicable for the whole of z. This ensures a finite enclosure when z contains the branch point $-1/e$.

Figure 2: Plot of the real part of $W_1(z)$ on the vertical segment $z = -1 + yi, |y| \leq 3$. The boxes show the range of the output intervals given input intervals $y = [a, a + h]$. The picture demonstrates continuity between the branch cut and the upper half plane: as intended, an imaginary part of $[-h, 0]$ (or $[-h/2, h/2]$, say, though not pictured here) in the input captures the jump discontinuity while $[0, h]$ does not. Where continuous, the output intervals converge nicely when $h \to 0$.

3   Bounds and series expansions

We proceed to state the inequalities needed for various error bounds in the algorithm.

3.1   Taylor series

Near the origin of the $k = 0$ branch, we have the Taylor series

$$W_0(z) = \sum_{n=1}^{\infty} \frac{(-n)^{n-1}}{n!} z^n.$$

Since $|n^{n-1}/n!| < e^n$, the truncation error on stopping before the $n = T$ term is bounded by $e^T |z|^T / (1 - e|z|)$ if $|z| < 1/e$.

3.2   Puiseux series

Near the branch point at $-1/e$ when $W(z) \approx -1$, the Lambert W function can be computed by means of a Puiseux series. This especially useful for intervals containing the point $-1/e$ itself, since we can compute a finite enclosure whereas enclosures based on $W'(z)$ blow up. If $\alpha = \sqrt{2(ez + 1)}$, then provided that $|\alpha| < \sqrt{2}$, we have

$$W_k(z) = \begin{cases} B(\alpha) & \text{if } k = 0 \\ B(-\alpha) & \text{if } k = -1 \text{ and } \operatorname{Im}(z) \geq 0 \\ B(-\alpha) & \text{if } k = +1 \text{ and } \operatorname{Im}(z) < 0 \end{cases}$$

7

where

$$B(\xi) = W\left(\frac{\xi^2 - 2}{2e}\right) = \sum_{n=0}^{\infty} c_n \xi^n. \tag{2}$$

Note that $W_{\pm 1}$ have one-sided branch cuts on $(-\infty, 0)$ and $(-1/e, 0)$. In the opposite upper and lower half planes, there is only a single cut on $(-\infty, 0)$ so the point $-1/e$ does not need to be treated specially.

In (2), the appropriate branches of W are implied so that $B(\xi)$ is analytic on $|\xi| < \sqrt{2}$. In terms of the standard branch cuts W_k, that is

$$k = \begin{cases} 0 & \text{if } -\pi/2 < \arg(\xi) \le \pi/2 \\ 1 & \text{if } \pi < \arg(\xi) < -\pi/2 \\ -1 & \text{otherwise.} \end{cases}$$

The coefficients c_n are rational numbers

$$c_0 = -1, \ c_1 = 1, \ c_2 = -\frac{1}{3}, \ c_3 = \frac{11}{72}, \ c_4 = -\frac{43}{540}, \ldots$$

which can be computed recursively. From singularity analysis, $|c_n| = O((1/\sqrt{2})^n)$, but we need an explicit numerical bound for computations. The following estimate is not optimal, but adequate for practical use.

Theorem 1. *The coefficients in (2) satisfy $|c_n| < 2 \cdot (4/5)^n$, or more simply, $|c_n| \le 1$.*

Proof. Numerical evaluation of W shows that $|2 + B(\xi)| < 2$ on the circle $|\xi| = 5/4$, so the Cauchy integral formula gives the result. $\square$

The verification can of course be done using interval arithmetic, as demonstrated in Figure 4. We stress that there is no circular dependency on Theorem 1 since the Puiseux series is not used for evaluation that far from from the branch point.

3.3 Asymptotic series

The Lambert W function has the asymptotic expansion

$$W_k(z) \sim L_1 - L_2 + \sum_{l=0}^{\infty} \sum_{m=1}^{\infty} c_{l,m} \sigma^l \tau^m \tag{3}$$

where

$$L_1 = \log(z) + 2\pi k i, \quad L_2 = \log(L_1), \quad \sigma = 1/L_1, \quad \tau = L_2/L_1 \tag{4}$$

and

$$c_{l,m} = \frac{(-1)^m}{m!} \begin{bmatrix} l+m \\ l+1 \end{bmatrix} \tag{5}$$

where $\begin{bmatrix} n \\ k \end{bmatrix}$ denotes an (unsigned) Stirling number of the first kind.

This expansion is valid for all k when $|z| \to \infty$, and also for $k \ne 0$ when $|z| \to 0$. In fact, (3) is not only an asymptotic series but (absolutely and uniformly) convergent for all sufficiently small $|\sigma|, |\tau|$. These properties of the expansion (3) were proved in [2].

The asymptotic behavior of the coefficients $c_{l,m}$ was studied further in [5], but that work did not give explicit inequalities. We will give an explicit bound for $|c_{l,m}|$, which permits us to compute $W_k(z)$ directly from (3) with a bound on the error in the relevant asymptotic regimes.

8

Lemma 2. *For all $n, k \geq 0$,*

$$\begin{bmatrix} n \\ k \end{bmatrix} \leq \frac{2^n n!}{k!}.$$

Proof. This follows by induction on the recurrence relation

$$\begin{bmatrix} n+1 \\ k \end{bmatrix} = n \begin{bmatrix} n \\ k \end{bmatrix} + \begin{bmatrix} n \\ k-1 \end{bmatrix}.$$

$\square$

Lemma 3. *For all $l, m \geq 0$, $|c_{l,m}| \leq 4^{l+m}$.*

Proof. By the previous lemma,

$$|c_{l,m}| \leq \frac{2^{l+m}(l+m)!}{(l+1)!m!} \leq 2^{l+m} \binom{l+m}{m} \leq 4^{l+m}.$$

$\square$

We can now restate (3) in the following effective form.

Theorem 4. *With σ, τ, L_1, L_2 defined as above, if $|\sigma| < 1/4$ and $|\tau| < 1/4$, and if $|z| > 1$ when $k = 0$, then*

$$W_k(z) = L_1 - L_2 + \sum_{l=0}^{L-1} \sum_{m=1}^{M-1} c_{l,m} \sigma^l \tau^m + \varepsilon_{L,M}(z)$$

with

$$|\varepsilon_{L,M}(z)| \leq \frac{4|\tau|(4|\sigma|)^L + (4|\tau|)^M}{(1-4|\sigma|)(1-4|\tau|)}.$$

Proof. Under the stated conditions, the series (3) converges to $W_k(z)$, by the analysis in [2]. We can bound the tail as

$$\left| \sum_{l=L}^{\infty} \sum_{m=M}^{\infty} c_{l,m} \sigma^l \tau^m \right| \leq \sum_{l=0}^{\infty} \sum_{m=M}^{\infty} (4|\sigma|)^l (4|\tau|)^m + \sum_{l=L}^{\infty} \sum_{m=1}^{\infty} (4|\sigma|)^l (4|\tau|)^m.$$

Evaluating the bivariate geometric series gives the result. $\square$

3.4 Bounds for the derivative

Finally, we give an rigorous global bound for the magnitude of W'. Since we want to compute W with small *relative* error, the estimate for $|W'(z)|$ should be optimal (up to a small constant factor) anywhere, including near singularities. We did not obtain a single neat expression that covers $W_k(z)$ adequately for all k and z, so a few case distinctions are made.

W' like W is a multivalued function, and whenever we fix a branch for W, we fix the corresponding branch for W'. Exactly on a branch cut, W' is therefore finite (except at a branch point) and equal to the directional derivative taken along the branch cut, so we must deal with the branch cut discontinuity separately when bounding perturbations in W if z crosses the cut.

The derivative of the Lambert W function can be written as

$$W'(z) = \frac{1}{(1+W(z))e^{W(z)}} = \frac{1}{z} \frac{W(z)}{1+W(z)}$$

where a limit needs to be taken in the rightmost expression for $W_0(z)$ near $z = 0$. The rightmost expression also shows that $W'(z) \approx 1/z$ when $|W(z)|$ is large. Bounding $|\mathrm{Im}(W_k(z))|$ from below gives the following.

Theorem 5. *For* $|k| \geq 2,$

$$|W_k'(z)| \leq \frac{1}{|z|}\frac{(2k-2)\pi}{(2k-2)\pi - 1} \leq \frac{1}{|z|}\frac{2\pi}{2\pi - 1} \leq \frac{1.2}{|z|}.$$

Also, if $k = 1$ *and* $\mathrm{Im}(z) \geq 0$, *or if* $k = -1$ *and* $\mathrm{Im}(z) < 0$, *then*

$$|W_k'(z)| \leq \frac{1}{|z|}\frac{\pi}{\pi - 1} \leq \frac{1.5}{|z|}.$$

For large $|z|$, the following two results are convenient.

Theorem 6. *If* $|z| > e$, *then for any* k,

$$|W_k'(z)| \leq \frac{1}{|z|}\frac{W_0(|z|)}{W_0(|z|) - 1}.$$

Proof. The inequality $|W_k(z)| \geq W_0(|z|)$ holds for all z (this is easily proved from the inverse function relationship defining W), giving the result. $\qquad\square$

Theorem 7. *If* $|z| \geq \left(\frac{1}{2} + (2|k|+1)\pi\right)e^{-1/2}$, *or more simply if* $|z| \geq 4(|k|+1)$, *then*

$$|W_k'(z)| \leq \frac{1}{|z|}.$$

Proof. Let $a = \mathrm{Re}(W_k(z))$. We have $|W_k(z)/(1 + W_k(z))| \leq 1$ when $a \geq -1/2$. If $a < -1/2$, then $|z| = |W_k(z)e^{W_k(z)}| < (|a| + (2|k|+1)\pi)e^a < (\frac{1}{2} + (2|k|+1)\pi)e^{-1/2}$. $\qquad\square$

It remains to bound $|W_k'(z)|$ for $k \in \{-1, 0, 1\}$ in the cases where z may be near the branch point at $-1/e$. This can be accomplished as follows.

Theorem 8. *For any* k,

$$|W_k'(z)| \leq \frac{1}{|z|}\max\left(3, \frac{1.5}{\sqrt{|ez+1|}}\right).$$

Proof. If $|W(z)+1| \geq 1/2$, then $|W(z)/(W(z)+1)| \leq 3$. Now consider the case $W(z)+1 = \varepsilon$ for some $|\varepsilon| \leq 1/2$. Then we must have $|ez+1| \leq |\varepsilon|^2$, due to the Taylor expansion

$$(-1+\varepsilon)e^{-1+\varepsilon} + e^{-1} = \frac{1}{e}\left(\frac{\varepsilon^2}{2} + \frac{\varepsilon^2}{3} + \frac{\varepsilon^2}{8} + \dots\right).$$

This implies that

$$\left|\frac{W(z)}{W(z)+1}\right| = \frac{|\varepsilon - 1|}{|\varepsilon|} \leq \frac{1 + |\varepsilon|}{|\varepsilon|} \leq \frac{1.5}{\sqrt{|ez+1|}}.$$

$\qquad\square$

Theorem 8 can be used practice, provided that we use a different bound when $k = 0$ and $z \approx 0$ (also, when $z \approx -1/e$ and $W_k(z) \not\approx -1$). However, it is worth making a few case distinctions and slightly complicating the formulas to tighten the error propagation for $k = -1, 0, 1$. For these branches, we implement the following inequalities.

Theorem 9. *Let* $t = |ez + 1|$.

1. *If* $|z| \leq 64$, *then*

$$|W_0'(z)| \leq \frac{2.25}{\sqrt{t(1+t)}}.$$

2. *If $|z| \geq 1$, then*

$$|W_0'(z)| \leq \frac{1}{|z|}.$$

3. *If $\mathrm{Re}(z) \geq 0$, or if $\mathrm{Im}(z) < 0$ when $k = -1$ (respectively $\mathrm{Im}(z) \geq 0$ when $k = 1$), then*

$$|W_{\pm 1}(z)| \leq \frac{1}{|z|}\left(1 + \frac{1}{4 + |z|^2}\right).$$

4. *For all z,*

$$|W_{\pm 1}(z)| \leq \frac{1}{|z|}\left(1 + \frac{23}{32}\frac{1}{\sqrt{t}}\right).$$

Proof. The inequalities can be verified by interval computations on a bounded region (since $1/|z|$ is an upper bound for sufficiently large $|z|$) excluding the neighborhoods of the branch points. These computations can be done by bootstrapping from Theorem 8. Close to $-1/e$, Theorem 1 applies, and an argument similar to that in Theorem 8 can be used close to 0. (We omit the straightforward but lengthy numerical details.) $\qquad\square$

It is clearly possible to make the bounds sharper, not least by adding more case distinctions, but these formulas are sufficient for our purposes, easy to implement, and cheap to evaluate. The implementation requires only the extraction of lower or upper bounds of intervals and unsigned floating-point operations with directed rounding (assuming that $ez + 1$ has been computed using interval arithmetic).

4 Alternative branch cuts

If the input z is an exact floating-point number, then we can always pinpoint its location in relation to the standard branch cuts of W. However, if the input is generated by an interval computation, it might look like $z = -10 + [\pm\varepsilon]i$ where the sign of $\mathrm{Im}(z)$ is ambiguous. If we want to compute solutions of $we^w = z$ in this case, the standard branches W_k do not work well because the jump discontinuity on the branch cut prevents the output intervals from converging when $\varepsilon \to 0$.

Likewise, when evaluating an integral or a solution of a differential equation involving W, say $\int_a^b f(z, W(g(z)))dz$, we might need to consider paths that would cross the standard branch cuts. We already saw an example with the application of the Cauchy integral formula to the Puiseux series coefficients in Section 3.2.

It is instructive to consider the treatment of square roots and logarithms, where the branch cut can be moved from $(-\infty, 0)$ to $(0, \infty)$ quite easily. The solutions of $w^2 = z$ are given by $w = \sqrt{z}, -\sqrt{z}$, but switching to $w = i\sqrt{-z}, -i\sqrt{-z}$ gives continuity along paths crossing the negative real axis. Similarly, for the solutions of $e^w = z$, we can switch from $w = \log(z) + 2k\pi i$ to $w = \log(-z) + (2k + 1)\pi i$.

The Lambert W function lacks a functional equation that simply would allow us to negate z. Instead, we define a set of alternative branches for W as follows:

- $W_{\mathrm{left}|k}(z)$ joins $W_k(z)$ for z in the upper half plane with $W_{k+1}(z)$ in the lower half plane, providing continuity to the left of the branch point at 0 (when $k \notin \{-1, 0\}$) or $-1/e$ (when $k \in \{-1, 0\}$). The branch cuts of this function thus extend from 0 or $-1/e$ to $+\infty$.

Figure 3: Plot of the real part (even function) and imaginary part (odd function) of $W(e^{\pi i \theta})$ with continuous analytic continuation on the Riemann surface of W. The branch used for evaluation is W_0 on $\theta \in [-0.5, 0.5]$, $W_{\text{left}|0}$ on $[0.5, 1.5]$, W_1 on $[1.5, 2.5]$, $W_{\text{left}|1}$ on $[2.5, 3.5]$, W_2 on $[3.5, 4.5]$, and $W_{\text{left}|2}$ on $[4.5, 5.5]$. Continuity is preserved whenever θ crosses an integer, that is, when $z = e^{\pi i \theta}$ crosses the real axis. The input intervals for θ have width $1/13$.

- $W_{\text{middle}}(z)$ joins $W_{-1}(z)$ in the upper half plane with $W_1(z)$ in the lower half plane, with continuity through the central segment $(-1/e, 0)$. This function extends the real analytic function $W_{-1}(x), x \in (-1/e, 0)$ to a complex analytic function on $z \in \mathbb{C} \setminus (-\infty, -1/e] \cup [0, \infty)$, unlike the standard branch $W_{-1}(z)$ where the real-valued segment lies precisely on the branch cut.

We follow the principle of counter-clockwise continuity to define the values of these alternative branches on their branch cuts (absent use of signed zero).

In the Arb implementation, the user can select the respective modified branch cuts by passing a special value in the *flags* field instead of the default value 0, namely

```
acb_lambertw(res, z, k, ACB_LAMBERTW_LEFT, prec)
acb_lambertw(res, z, k, ACB_LAMBERTW_MIDDLE, prec)
```

where $k = -1$ should be set in the second case.

We implement the alternative branch cuts by splitting the input into $z_a = \text{Re}(z) + (\text{Im}(z) \cap [0, \infty))i$ and $z_b = \text{Re}(z) + (-\text{Im}(z) \cap [0, \infty))i$. If the standard branches to be taken below and above the cut have index k and k' respectively, then we compute $W(z)$ as $W_k(z_a) \cup \overline{W_{-k'}(z_b)}$. Conjugation is used to get a continuous evaluation of $W_{k'}(\text{Re}(z) + (\text{Im}(z) \cap (-\infty, 0))i)$, in light of our convention to work with closed intervals and make the standard branches W_k continuous from above on the cut.

We observe that for $W_{\text{middle}}(z)$ the Puiseux expansion at $-1/e$ is valid in all directions, as is the asymptotic expansion at 0 with $L_1 = \log(-z)$ and $L_2 = \log(-L_1)$. Further, $W_{\text{left}|k}(z)$ is given by the asymptotic expansion with $L_1 = \log(-z) + (2k+1)\pi i, L_2 = \log(L_1)$ when $|z| \to \infty$. These formulas could be used directly instead of case splitting where applicable.

12

Figure 4: Plot of the real and imaginary part and the absolute value (black) of $2 + W((z^2 - 2)/(2e))$, $z = 1.25e^{\pi i\theta}$ with continuous analytic continuation. The function argument $(z^2 - 2)/(2e)$ traces two loops around the branch point at $-1/e$, passing through the branches $0, +1, -1$, and back to 0. From left to right, the branch used for evaluation cycles through $W_0, W_{\text{left}|0}, W_1, W_{\text{middle}}, W_{-1}, W_{\text{left}|-1}, W_0$. Input intervals have been subdivided adaptively to show the absolute value bound.

5   Testing and benchmarks

z	10	100	1000	10000
10	3.36	7.12	1.60	1.50
10^{10}	3.64	6.92	1.65	1.53
$10^{10^{20}}$	3.46	8.39	1.91	1.67
$10i$	13.20	8.68	4.71	3.27
$-10^{10^{20}}$	3.69	29.75	7.53	4.59
$-1/e + 10^{-100}$	4.57	2.33	2.23	1.97
$-1/e - 10^{-100}$	4.43	2.36	7.08	2.89

Table 1: Time to compute $w = W_0(z)$, relative to the time to compute e^w, at a precision of 10, 100, 1000 and 10000 decimal digits.

We have tested the implementation in Arb in various ways, most importantly to verify that correct inclusions are being computed, but also to make sure that output intervals are reasonably tight.

The automatic unit test included with the library generates overlapping random input intervals z_1, z_2 (sometimes placed very close to $-1/e$), computes $w_1 = W_k(z_1)$ and $w_2 = W_k(z_2)$ at different levels of precision (sometimes directly invoking the asymptotic expansion with a random number of terms instead of calling the main Lambert W function implementation), checks that the intervals w_1 and w_2 overlap, and also checks that $w_1 e^{w_1}$ contains z_1. The conjugate symmetry is also tested. These checks give a strong test of correctness.

We have also done separate tests to verify that the error bounds converge for exact floating-point input when the precision is increased, and further ad hoc tests have been done

13

to test a variety of easy and difficult cases at different precisions.

At low precision, the absolute time to evaluate W for a "normal" input z is about 10^{-6} seconds when W is real and 10^{-5} seconds when W is complex (on an Intel i5-4300U CPU). For instance, creating a 1000 by 1000 pixel domain coloring plot of $W_0(z)$ on $[-5, 5] + [-5, 5]i$ takes 12 seconds.

Table 1 shows normalized timings for `acb_lambertw`. The higher relative overhead when W is complex mainly results from less optimized precision handling in the floating-point code (which could be improved in a future version), together with some extra overhead for the branch test.

We show the output (converted to decimal intervals using `arb_printn`) for a few of the test cases in the benchmark. For $z = 10$, the following results are computed at the respective levels of precision:

```
[1.745528003 +/- 3.82e-10]
[1.7455280027{...79 digits...}0778883075 +/- 4.71e-100]
[1.7455280027{...979 digits...}5792011195 +/- 1.97e-1000]
[1.7455280027{...9979 digits...}9321568319 +/- 2.85e-10000]
```

For $z = 10^{10^{20}}$, we get:

```
[2.302585093e+20 +/- 3.17e+10]
[2302585092994045683554.9134111633{...59 digits...}5760752900 +/- 6.06e-80]
[2302585092994045683554.9134111633{...959 digits...}8346041370 +/- 3.55e-980]
[2302585092994045683554.9134111633{...9959 digits...}2380817535 +/- 6.35e-9980]
```

For $z = -1/e + 10^{-100}$, the input interval overlaps with the branch point at 10 and 100 digits, showing a potential small imaginary part in the output, but at higher precision the imaginary part disappears:

```
[-1.000 +/- 3.18e-5] + [+/- 2.79e-5]i
[-1.0000000000{...28 digits...}0000000000 +/- 3.81e-50] + [+/- 2.76e-50]i
[-0.9999999999{...929 digits...}9899904389 +/- 2.99e-950]
[-0.9999999999{...9929 digits...}9452369126 +/- 5.45e-9950]
```

6 Automatic differentiation

Finally, we consider the computation of derivatives $W^{(n)}$, or more generally $(W \circ f)^{(n)}$ for an arbitrary function f. That is, given a power series $f \in \mathbb{C}[[x]]$, we want to compute the power series $W(f)$ truncated to length $n + 1$.

The higher derivatives of W can be calculated using recurrence relations as discussed in [2], but it is more efficient to use formal Newton iteration in the ring $\mathbb{C}[[x]]$ to solve the equation $we^w = f$. That is, given a power series w_j correct to n terms, we compute

$$w_{j+1} = w_j - \frac{w_j e^{w_j} - f}{e^{w_j} + w_j e^{w_j}}$$

which is correct to $2n$ terms.

Indeed, this approach allows us to compute the first n derivatives of W or $W \circ f$ (when the first n derivatives of f are given) in $O(\mathsf{M}(n))$ operations where $\mathsf{M}(n)$ is the complexity of polynomial multiplication. With FFT based multiplication, we have $\mathsf{M}(n) = O(n \log n)$.

This method is implemented by the Arb functions `arb_poly_lambertw_series` (for real polynomials) and `acb_poly_lambertw_series` (for complex polynomials).

Since the low n coefficients of w_{j+1} and w_j are identical mathematically, we simply copy these coefficients instead of performing the full subtraction (avoiding needless inflation of the enclosures). A further important optimization in this algorithm is to save the constant

term $e_0 = [x^0]e^w$ so that e^{w_j} can be computed as $e_0 e^{w_j - [x^0]w_j}$. This avoids a transcendental function evaluation, which is expensive and moreover can be ill-conditioned, leading to greatly inflated enclosures. The performance could be improved further by a constant factor by saving the partial Newton iterations done internally for power series division and exponentials.

Empirically, the Newton iteration scheme is reasonably numerically stable, permitting the evaluation of high order derivatives with modest extra precision even in interval arithmetic. For example, computing 10000 terms in the series expansion of $h(x) = W_0(e^{1+x})$ at 256-bit precision takes 2.8 seconds, giving $[x^{10000}]h(x)$ as

```
[-6.02283194399026390e-5717 +/- 5.56e-5735].
```

7 Discussion

A number of improvements could be pursued in future work.

The algorithm presented here is correct in the sense that it computes a validated enclosure for $W_k(z)$, absent any bugs in the code. It is also easy to see that the enclosures converge when the input intervals converge and the precision is increased accordingly (as long as a branch cut is not crossed), under the assumption that the floating-point approximation is computed accurately. However, we have made no attempt to prove that the floating-point approximation is computed accurately beyond the usual heuristic reasoning and experimental testing.

Although the focus is on interval arithmetic, we note that applying Ziv's strategy [9] allows us to compute floating-point approximations of $W_k(z)$ with certified correct rounding. This requires only a simple wrapper around the interval implementation without the need for separate analysis of floating-point rounding errors. A rigorous floating-point error analysis for computing the Lambert W function without the use of interval arithmetic seems feasible, certainly for real variables but probably also for complex variables.

We use a first order bound based on $|W'(z)|$ for error propagation when z is inexact. For wide z, more accurate bounds could be achieved using higher-order estimates. Simple and tight bounds for $|W^{(n)}(z)|$ for small n would be a useful addition.

For very wide intervals z, optimal enclosures could be determined by evaluating W at two or more points to find the extreme values. This is most easily done in the real case, but suitable monotonicity conditions could be determined for complex variables as well.

The implementation in Arb is designed for arbitrary precision. For low precision, the main approximation is usually computed using `double` arithmetic, but the certification uses arbitrary-precision arithmetic which consumes the bulk of the time. Using validated `double` or double-double arithmetic for the certification would be significantly faster.

References

[1] F. Chapeau-Blondeau and A. Monir. Numerical evaluation of the Lambert W function and application to generation of generalized Gaussian noise with exponent 1/2. *IEEE Transactions on Signal Processing*, 50(9):2160–2165, 2002.

[2] R. M. Corless, G. H. Gonnet, D. E. G. Hare, D. J. Jeffrey, and D. E. Knuth. On the Lambert W function. *Advances in Computational Mathematics*, 5(1):329–359, 1996.

[3] F. Johansson. Efficient implementation of elementary functions in the medium-precision range. In *22nd IEEE Symposium on Computer Arithmetic*, ARITH22, pages 83–89, 2015.

[4] F. Johansson. Arb: Efficient arbitrary-precision midpoint-radius interval arithmetic. *IEEE Transactions on Computers*, PP(99):1–1, 2017. `http://dx.doi.org/10.1109/TC.2017.2690633` (to appear).

[5] G. A. Kalugin and D. J. Jeffrey. Convergence in C of series for the Lambert W function. *arXiv preprint arXiv:1208.0754*, 2012.

[6] P. W. Lawrence, R. M. Corless, and D. J. Jeffrey. Algorithm 917: complex double-precision evaluation of the Wright ω function. *ACM Transactions on Mathematical Software*, 38(3):20, 2012.

[7] R. E. Moore. *Methods and applications of interval analysis.* SIAM, 1979.

[8] D. Veberič. Lambert W function for applications in physics. *Computer Physics Communications*, 183(12):2622–2628, 2012.

[9] A. Ziv. Fast evaluation of elementary mathematical functions with correctly rounded last bit. *ACM Transactions on Mathematical Software*, 17(3):410–423, 1991.